

```
import numpy as np
import pandas as pd

df = pd.read_csv('/spam.csv')

df.head(5)
```

| | v1 | v2 | Unnamed: 2 | Unnamed: 3 | Unnamed: 4 |
|---|------|---|------------|------------|------------|
| 0 | ham | Go until jurong point, crazy.. Available only ... | NaN | NaN | NaN |
| 1 | ham | Ok lar... Joking wif u oni... | NaN | NaN | NaN |
| 2 | spam | Free entry in 2 a wkly comp to win FA Cup fina... | NaN | NaN | NaN |
| 3 | ham | U dun say so early hor... U c already then say... | NaN | NaN | NaN |

```
df.shape

↳ (5572, 5)
```

Steps Involved in this Process :

1)Data Cleaning 2)EDA 3)Text Preprocessing 4)ModelBuilding 5)Evaluation 6)Improvement

1.Data Cleaning

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    v1          5572 non-null   object
1    v2          5572 non-null   object
2    Unnamed: 2   50 non-null     object
3    Unnamed: 3   12 non-null     object
4    Unnamed: 4    6 non-null     object
dtypes: object(5)
memory usage: 217.8+ KB

# drop last 3 cols
df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],inplace=True)

df.head(5)
```

| | v1 | v2 |
|---|------|---|
| 0 | ham | Go until jurong point, crazy.. Available only ... |
| 1 | ham | Ok lar... Joking wif u oni... |
| 2 | spam | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham | U dun say so early hor... U c already then say... |
| 4 | ham | Nah I don't think he goes to usf, he lives aro... |

```
# renaming the cols
df.rename(columns={'v1':'target','v2':'text'},inplace=True)
df.sample(5)
```

| | target | text |
|------|--------|---|
| 3133 | ham | Ok going to sleep. Hope i can meet her. |
| 5015 | ham | Hey gals...U all wanna meet 4 dinner at nite? |
| 3465 | ham | Actually fuck that, just do whatever, do find ... |
| 2397 | ham | Neshanth..tel me who r u? |
| 4212 | ham | I attended but nothing is there. |

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
```

```
df['target'] = encoder.fit_transform(df['target'])
```

```
df.head()
```

| | target | text |
|---|--------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... |
| 1 | 0 | Ok lar... Joking wif u oni... |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... |

```
data = df['text'].copy()
```

```
# missing values
df.isnull().sum()
```

```
target    0
text      0
dtype: int64
```

```
# check for duplicate values
df.duplicated().sum()
```

```
409
```

```
# remove duplicates
df = df.drop_duplicates(keep='first')
```

```
df.duplicated().sum()
```

```
0
```

```
df.shape
```

```
(5163, 2)
```

```
df.head()
```

| | target | text |
|---|--------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... |
| 1 | 0 | Ok lar... Joking wif u oni... |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... |


```
data = df['text'].copy()
```

```
*2)EDA *
```

```
df['target'].value_counts()
```

```
0    4516
1     647
Name: target, dtype: int64
```

```
import matplotlib.pyplot as plt
plt.pie(df['target'].value_counts(), labels=['ham', 'spam'], autopct="%0.2f")
plt.show()
```



Data is Imbalanced

```
import nltk
!pip install nltk
nltk.download('punkt')

Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk) (8.1.6)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk) (1.3.1)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk) (2022.10.31)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk) (4.65.0)
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
True
```

```
df['num_characters'] = df['text'].apply(len)
```

```
df.head()
```

| | target | text | num_characters |
|---|--------|---|----------------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 |

```
df['num_words'] = df['text'].apply(lambda x : len(nltk.word_tokenize(x)))
```

```
df.head()
```

| | target | text | num_characters | num_words |
|---|--------|---|----------------|-----------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 |

```
df['num_sentences'] = df['text'].apply(lambda x : len(nltk.sent_tokenize(x)))
```

```
df.head()
```

| | target | text | num_characters | num_words | num_sentences |
|---|--------|---|----------------|-----------|---------------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 | 1 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 | 1 |

```
df[['num_characters', 'num_words', 'num_sentences']].describe()
```

| | num_characters | num_words | num_sentences |
|--------------|----------------|-------------|---------------|
| count | 5163.000000 | 5163.000000 | 5163.000000 |
| mean | 78.811156 | 18.444122 | 1.963974 |
| std | 58.153491 | 13.327940 | 1.448008 |
| min | 2.000000 | 1.000000 | 1.000000 |
| 25% | 36.000000 | 9.000000 | 1.000000 |
| 50% | 60.000000 | 15.000000 | 1.000000 |
| 75% | 116.000000 | 26.000000 | 2.000000 |
| max | 910.000000 | 220.000000 | 38.000000 |

```
#ham
```

```
df[df['target'] == 0][['num_characters', 'num_words', 'num_sentences']].describe()
```

| | num_characters | num_words | num_sentences |
|--------------|----------------|-------------|---------------|
| count | 4516.000000 | 4516.000000 | 4516.000000 |
| mean | 70.428919 | 17.123782 | 1.820195 |
| std | 56.356310 | 13.493970 | 1.383657 |
| min | 2.000000 | 1.000000 | 1.000000 |
| 25% | 34.000000 | 8.000000 | 1.000000 |
| 50% | 52.000000 | 13.000000 | 1.000000 |
| 75% | 90.000000 | 22.000000 | 2.000000 |
| max | 910.000000 | 220.000000 | 38.000000 |

```
#spam
```

```
df[df['target'] == 1][['num_characters', 'num_words', 'num_sentences']].describe()
```

| | num_characters | num_words | num_sentences |
|--------------|----------------|------------|---------------|
| count | 647.000000 | 647.000000 | 647.000000 |
| mean | 137.318393 | 27.659969 | 2.967543 |
| std | 30.096861 | 7.038065 | 1.490069 |
| min | 13.000000 | 2.000000 | 1.000000 |
| 25% | 130.000000 | 25.000000 | 2.000000 |
| 50% | 148.000000 | 29.000000 | 3.000000 |
| 75% | 157.000000 | 32.000000 | 4.000000 |
| max | 223.000000 | 46.000000 | 9.000000 |

```
import seaborn as sns
```

```
plt.figure(figsize=(12,6))
```

```
sns.histplot(df[df['target']==0]['num_characters'])
```


```
sns.histplot(df[df['target']==1]['num_characters'],color='red')
```

<Axes: xlabel='num_characters', ylabel='Count'>



```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target']==0]['num_words'])
sns.histplot(df[df['target']==1]['num_words'],color='red')
```

<Axes: xlabel='num_words', ylabel='Count'>



```
sns.pairplot(df,hue='target')
```

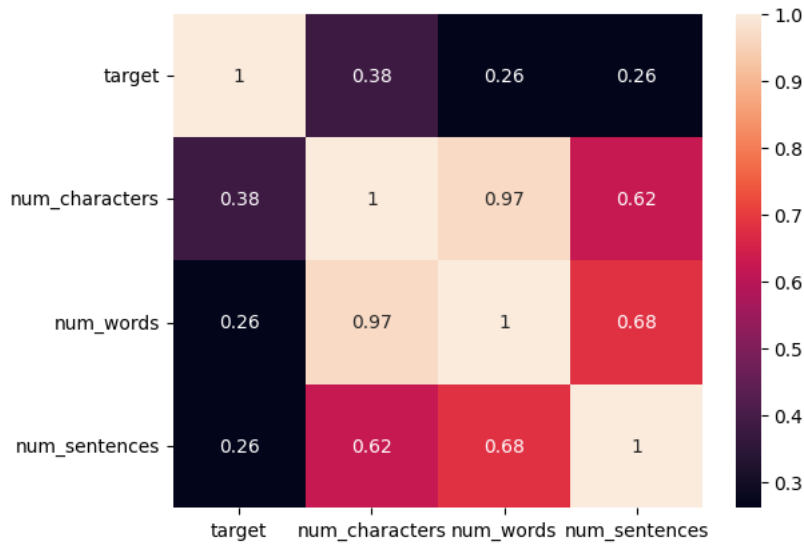
```
<seaborn.axisgrid.PairGrid at 0x7fec31339840>
```



```
sns.heatmap(df.corr(),annot=True)
```

```
<ipython-input-35-8df7bcac526d>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is
sns.heatmap(df.corr(),annot=True)
```

```
<Axes: >
```



3.Data Preprocessing

Steps involved in this process are :

- Lower case
- Tokenization
- Removing special characters
- Removing stop words and punctuation
- Stemming

```
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
stopwords.words('english')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

```
['i',
 'me',
 'my',
 'myself',
 'we',
 'our',
 'ours',
 'ourselves',
 'you',
 "you're",
 "you've",
 "you'll",
 "you'd",
 'your',
 'yours',
 'yourself',
 'yourselves',
 'he',
 'him',
 'his',
 'himself',
 'she',
 "she's",
 'her',
 'hers',
 'herself',
 'it',
 "it's",
 'its',
 'itself',
```

```

    'they',
    'them',
    'their',
    'theirs',
    'themselves',
    'what',
    'which',
    'who',
    'whom',
    'this',
    'that',
    "that'll",
    'these',
    'those',
    'am',
    'is',
    'are',
    'was',
    'were',
    'be',
    'been',
    'being',
    'have',
    'has',
    'had',
    'having',

import string
string.punctuation

'!"#$%&\'()*+,-./:;<=>?@[\\]^_`{|}~'

from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()

def transform_text(text) :
    text = text.lower()
    text = nltk.word_tokenize(text)
    y = []

    for i in text :
        if i.isalnum():
            y.append(i)

    text = y[:]
    y.clear()

    for i in text:
        if i not in stopwords.words('english') and i not in string.punctuation :
            y.append(i)

    text = y[:]
    y.clear()

    for i in text :
        y.append(ps.stem(i))

    return " ".join(y)

transform_text('Hi How are You 20')

'hi 20'

df['text'][2000]

'But i'll b going 2 sch on mon. My sis need 2 take smth.'


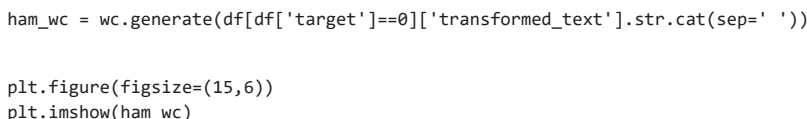
transform_text(df['text'][2000])

'b go 2 sch mon si need 2 take smth'

df['transformed_text'] = df['text'].apply(transform_text)
df.head()

```

```
from wordcloud import WordCloud
wc = WordCloud( width=500,height=500,min_font_size=10,background_color='white')
spam_wc = wc.generate(df[df['target']==1]['transformed_text'].str.cat(sep=' '))
plt.figure(figsize=(15,6))
plt.imshow(spam_wc)
```




```
for word in msg.split():
    spam_corpus.append(word)

len(spam_corpus)

9861

ham_corpus = []
for msg in df[df['target']==0]['transformed_text'].tolist() :
    for word in msg.split():
        ham_corpus.append(word)
```

```
len(ham_corpus )

35306
```

```
from collections import Counter
pd.DataFrame(Counter(ham_corpus).most_common(10))
```

| | 0 | 1 |
|---|------|-----|
| 0 | u | 883 |
| 1 | go | 404 |
| 2 | get | 349 |
| 3 | gt | 288 |
| 4 | lt | 287 |
| 5 | 2 | 284 |
| 6 | come | 275 |
| 7 | got | 236 |
| 8 | know | 236 |
| 9 | like | 234 |

```
# Text Vectorization
# using Bag of Words
df.head()
```

| | target | text | num_characters | num_words | num_sentences | transformed_text |
|---|--------|---|----------------|-----------|---------------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 | go jurong point crazi avail bugi n great world... |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 | ok lar joke wif u oni |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 | free entri 2 wkli comp win fa cup final tkt 21... |
| 3 | 0 | U dun say so early hor... U c already then sav... | 49 | 13 | 1 | u dun say earli hor u c already sav |

4.Model Building

```
from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
cv = CountVectorizer()
tfidf = TfidfVectorizer(max_features=3000)

x = tfidf.fit_transform(df['transformed_text']).toarray()

x.shape

(5163, 3000)

y = df['target'].values

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=2)

from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB
from sklearn.metrics import accuracy_score,confusion_matrix,precision_score
gnb = GaussianNB()
```

```

mnb = MultinomialNB()
bnb = BernoulliNB()

gnb.fit(x_train,y_train)
y_pred1 = gnb.predict(x_test)
print(accuracy_score(y_test,y_pred1))
print(precision_score(y_test,y_pred1))
print(confusion_matrix(y_test,y_pred1))

0.861568247821878
0.452991452991453
[[784 128]
 [ 15 106]]

mnb.fit(x_train,y_train)
y_pred2 = mnb.predict(x_test)
print(accuracy_score(y_test,y_pred2))
print(precision_score(y_test,y_pred2))
print(confusion_matrix(y_test,y_pred2))

0.9719264278799613
0.9893617021276596
[[911  1]
 [ 28  93]]

bnb.fit(x_train,y_train)
y_pred3 = bnb.predict(x_test)
print(accuracy_score(y_test,y_pred3))
print(precision_score(y_test,y_pred3))
print(confusion_matrix(y_test,y_pred3))

0.978702807357212
0.9714285714285714
[[909  3]
 [ 19 102]]

# tfidf --> MNB

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

svc = SVC(kernel='sigmoid', gamma=1.0)
knc = KNeighborsClassifier()
mnb = MultinomialNB()
dtc = DecisionTreeClassifier(max_depth=5)
lrc = LogisticRegression(solver='liblinear', penalty='l1')
rfc = RandomForestClassifier(n_estimators=50, random_state=2)
abc = AdaBoostClassifier(n_estimators=50, random_state=2)
bc = BaggingClassifier(n_estimators=50, random_state=2)
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
xgb = XGBClassifier(n_estimators=50, random_state=2)

clfs = {
    'SVC': svc,
    'KN': knc,
    'NB': mnb,
    'DT': dtc,
    'LR': lrc,
    'RF': rfc,
    'AdaBoost': abc,
    'BgC': bc,
    'ETC': etc,
    'GBDT': gbdt,
    'xgb': xgb
}

def train_classifier(clf,x_train,y_train,X_test,y_test):
    clf.fit(x_train,y_train)

```

```

y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test,y_pred)
precision = precision_score(y_test,y_pred)

return accuracy,precision

train_classifier(svc,x_train,y_train,x_test,y_test)

(0.978702807357212, 0.9805825242718447)

accuracy_scores = []
precision_scores = []

for name,clf in clfs.items():

    current_accuracy,current_precision = train_classifier(clf, x_train,y_train,x_test,y_test)

    print("For ",name)
    print("Accuracy - ",current_accuracy)
    print("Precision - ",current_precision)

    accuracy_scores.append(current_accuracy)
    precision_scores.append(current_precision)

For SVC
Accuracy - 0.978702807357212
Precision - 0.9805825242718447
For KN
Accuracy - 0.9225556631171346
Precision - 1.0
For NB
Accuracy - 0.9719264278799613
Precision - 0.9893617021276596
For DT
Accuracy - 0.936108422071636
Precision - 0.8021978021978022
For LR
Accuracy - 0.9554695062923524
Precision - 0.9310344827586207
For RF
Accuracy - 0.972894482090997
Precision - 0.979381443298969
For AdaBoost
Accuracy - 0.9622458857696031
Precision - 0.9270833333333334
For BgC
Accuracy - 0.9641819941916747
Precision - 0.868421052631579
For ETC
Accuracy - 0.9777347531461762
Precision - 0.9803921568627451
For GBDT
Accuracy - 0.9535333978702807
Precision - 0.9397590361445783
For xgb
Accuracy - 0.9622458857696031
Precision - 0.91

performance_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':accuracy_scores,'Precision':precision_scores}).sort_values('Precision',
performance_df

```

| | Algorithm | Accuracy | Precision |
|----|-----------|----------|-----------|
| 1 | KN | 0.922556 | 1.000000 |
| 2 | NB | 0.971926 | 0.989362 |
| 0 | SVC | 0.978703 | 0.980583 |
| 8 | ETC | 0.977735 | 0.980392 |
| 5 | RF | 0.972894 | 0.979381 |
| 9 | GBDT | 0.953533 | 0.939759 |
| 4 | LR | 0.955470 | 0.931034 |
| 6 | AdaBoost | 0.962246 | 0.927083 |
| 10 | xgb | 0.962246 | 0.910000 |
| 7 | BgC | 0.964182 | 0.868421 |
| 3 | DT | 0.936108 | 0.802198 |

```
n=67
input = x[n]
input_reshaped=input.reshape(1,-1)
prediction=mnf.predict(input_reshaped)
if prediction == 1 :
    print(data[n])
    print('SPAM')
else:
    print(data[n])
    print('NOT SPAM')
```

Urgent UR awarded a complimentary trip to EuroDisinc Trav, Aco&Entry41 Or 1000. To claim txt DIS to 87121 18+6*1.50(moreFrmMob. SPAM